

# IE 678 Deep Learning

## 02 – Feedforward Neural Networks

### **Part 1: Embeddings**

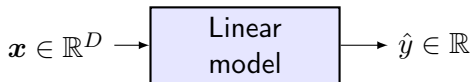
Prof. Dr. Rainer Gemulla

Universität Mannheim

Version: 2024-1

# From linear models to FNNs (1)


- Consider: prediction task with inputs  $x \in \mathcal{X}$  and outputs  $y \in \mathcal{Y}$ 
  - ▶ Goal: learn a function from  $\mathcal{X}$  to  $\mathcal{Y}$
- Simple approach: use a (generalized) **linear model**
  - ▶ Inputs must be real-valued feature vectors  $x \in \mathbb{R}^D$
  - ▶ Outputs are a real value (e.g., linear or logistic regression)
- Visually:




- Recall:  $\hat{y} = \phi(\mathbf{w}^\top \mathbf{x} + b)$ , where
  - ▶  $\mathbf{w} \in \mathbb{R}^D$  is a **weight vector** (one weight per feature, learned)
  - ▶  $b \in \mathbb{R}$  is a **bias term** (learned)
  - ▶  $\phi$  is a **mean function** (e.g., identity or logistic function)
- Problem: low representational capacity due to linearity assumption

## Example: Logistic regression (from ML course)

Data ( $\mathbf{x} \in \mathbb{R}^2$ )




Prediction ( $\hat{y} \in [0, 1]$ )



## From linear models to FNNs (2)

- Representational capacity can be addressed by feature engineering or using kernel methods
  - ▶ Allows to use arbitrary inputs spaces  $x \in \mathcal{X}$  by mapping them to real-valued vectors
  - ▶ To do so, uses *pre-specified* feature extractor  $f : \mathcal{X} \rightarrow \mathbb{R}^F$

- Visually:

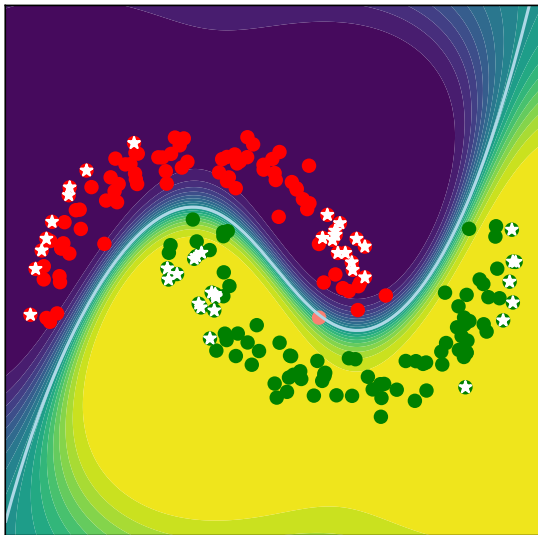


- Problem: which feature extractor?
  - ▶ Key to good performance
  - ▶ Hard to get right (domain experts, extensive experimentation, ...)
  - ▶ To see this: can you write a suitable feature extractor for classifying images? (if not, see [here](#))

## Example: L1VM (from ML course)

L1VM, RBF kernel, logistic regression


$$\lambda = 0.1, \sigma^2 = 0.571$$



## From linear models to FNNs (3)

- DL methods can be interpreted as an approach to *learn* features
  - ▶ Input objects  $x \in \mathcal{X}$  are transformed into dense, continuous, low-dimensional representations called **embeddings**  $z \in \mathbb{R}^Z$
  - ▶ Useful to represent complex objects (categorical data, textual data, graph data, tabular data, images, ...)
  - ▶ Think: complex to work with objects, simple to work with embeddings
  - ▶ Useful **embedding space** = goal of **representation learning**
  - ▶  $E$  = embedding dimensionality

- Visually:






- Key point: instead of engineering features manually, embeddings are learned from data → Main topic of this course
  - ▶ Embeddings also called: **latent code**, **distributed representations**
  - ▶ Embedding space also called: **latent space**

# Example: Document embeddings


804414 newswire stories, inputs = per-document rel. frequencies of 2000 most common word stems ( $x \in \mathbb{R}^{2000}$ ), shown here is 2D embedding ( $z \in \mathbb{R}^2$ ) of two different methods (left: linear, right: autoencoder)

**Fig. 4.** (A) The fraction of retrieved documents in the same class as the query when a query document from the test set is used to retrieve other test set documents, averaged over all 402,207 possible queries. (B) The codes produced by two-dimensional LSA. (C) The codes produced by a 2000-500-250-125-2 autoencoder.



# Encoders and prediction heads

- Functions that transform objects  $x \in \mathcal{X}$  to embeddings  $z \in \mathbb{R}^Z$  are known as **encoders**
- Functions that transform embeddings  $z \in \mathbb{R}^Z$  to predictions  $y \in \mathbb{R}$  are known as **prediction heads**
  - ▶ Can be linear or more complex
  - ▶ Typically much simpler than encoder
  - ▶ E.g., when prediction head is logistic regression, then positive and negative instances are ideally linearly separable in embedding space
- Visually:



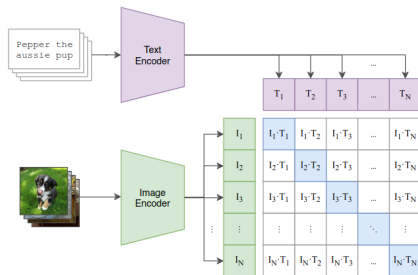
- Both encoder and prediction head are learned neural (sub)networks



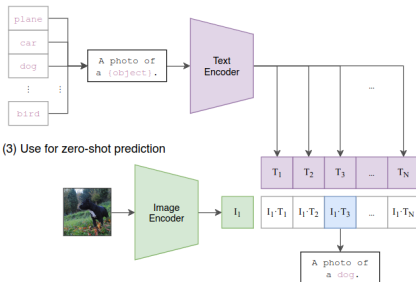
# Contrastive learning

- Embeddings can be used in other ways as well
- E.g., to compare objects, potentially across multiple modalities
  - ▶ Useful, for example, for zero- and few-shot prediction
  - ▶ Learned via a “contrastive learning” approach (more later)
- Example: [CLIP embeddings](#) for images and text

(1) Contrastive pre-training



(2) Create dataset classifier from label text




(3) Use for zero-shot prediction

Figure 1. Summary of our approach. While standard image models jointly train an image feature extractor and a linear classifier to predict some label, CLIP jointly trains an image encoder and a text encoder to predict the correct pairings of a batch of (image, text) training examples. At test time the learned text encoder synthesizes a zero-shot linear classifier by embedding the names or descriptions of the target dataset's classes.

# Structured prediction / deep generative models

- To handle more complex output spaces  $\mathcal{Y}$ , we may replace the prediction head by a component that “generates” output
- Functions that transform embeddings  $z \in \mathbb{R}^Z$  to (complex) outputs  $y \in \mathcal{Y}$  are known as **decoders**
  - ▶ Note: in such models, embedding dimensionality  $Z$  may or may not depend on input  $x$
- Visually:



## Example: unCLIP (DALL-E 2)




Figure 2: A high-level overview of unCLIP. Above the dotted line, we depict the CLIP training process, through which we learn a joint representation space for text and images. Below the dotted line, we depict our text-to-image generation process: a CLIP text embedding is first fed to an autoregressive or diffusion prior to produce an image embedding, and then this embedding is used to condition a diffusion decoder which produces a final image. Note that the CLIP model is frozen during training of the prior and decoder.