

IE 678 Deep Learning

02 – Feedforward Neural Networks

Part 2: Feedforward Neural Networks


Prof. Dr. Rainer Gemulla

Universität Mannheim

Version: 2024-1

Artificial neuron


- An **artificial neuron** (AN) is a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$
 - ▶ Inputs are vectors $\mathbf{x} \in \mathbb{R}^n$
 - ▶ Output is a value $y \in \mathbb{R}$, called **activation**
- f is taken from a family of functions that is parameterized by
 - ▶ A **weight vector** $\mathbf{w} \in \mathbb{R}^n$ (one weight per input)
 - ▶ A **bias** $b \in \mathbb{R}$
 - ▶ A **transfer function** or **activation function** $\phi : \mathbb{R} \rightarrow \mathbb{R}$
- Basic structure: $y = \phi(\mathbf{w}^\top \mathbf{x} + b)$
 - ▶ Computes the weighted sum $s = \mathbf{w}^\top \mathbf{x} + b$ of its inputs and bias
 - ▶ Passes s through the transfer function to obtain output y
 - ▶ ϕ can be deterministic or stochastic
- As before: bias can be replaced by an additional input $x_0 = 1$ and corresponding weight $w_0 = b$




Types of artificial neurons


- The **type** of an AN is determined by its transfer function ϕ
- An AN of a given type can represent the family of functions

$$F_{\phi} = \left\{ \mathbf{x} \rightarrow \phi(\mathbf{w}^{\top} \mathbf{x} + b) \mid \mathbf{w} \in \mathbb{R}^n, b \in \mathbb{R} \right\}$$

- Each function in this family can be represented by its bias and weight vector
- We will see later that types are usually specified up-front, whereas weights are learned
- The simplest type of neuron is a **constant neuron**
 - ▶ No inputs; output fixed value $x \in \mathbb{R}$
 - ▶ Notation (from now on): 


Example: Linear neuron / identity


- Uses $\phi(s) = s$
- Notation: 



- Simple but computationally limited
- We often but not always want non-linear transfer functions

Example: Logistic neuron


- Use logistic function $\phi(s) = \sigma(s) \stackrel{\text{def}}{=} \frac{1}{1+\exp(-s)}$
- Notation: 



- Gives a real-valued output that is smooth and bounded in $[0, 1]$
 - ▶ Negative activations mapped to value < 0.5
 - ▶ 0 activation mapped to 0.5
 - ▶ Positive activation mapped to value > 0.5
- Non-linear

Example: Stochastic binary neuron

- Also use logistic function
- But output of the logistic function is treated as a probability of producing a spike (1)
- I.e., $\phi(s) = \begin{cases} 1 & \text{with probability } \sigma(s) \\ 0 & \text{otherwise} \end{cases}$




- Defines a probability distribution over outputs
- Other neurons can also be made stochastic

What is an artificial neural network?

- A network of artificial neurons
 - ▶ A set of (artificial) neurons
 - ▶ Connections between neurons (directed or undirected)
- Many different **architectures**
 - ▶ How many neurons? Of which type?
 - ▶ Are there **output neurons**?
 - ▶ Are there **hidden neurons** (neither input nor output)?
 - ▶ Which neurons are connected?
 - ▶ Are connections directed or undirected?
 - ▶ Are there cycles?
- Picking the right architecture for the problem at hand is important and requires skill/thought/compute power
→ **Architecture engineering**
- Can represent a wide range of functions (*universal approximation theorem*)

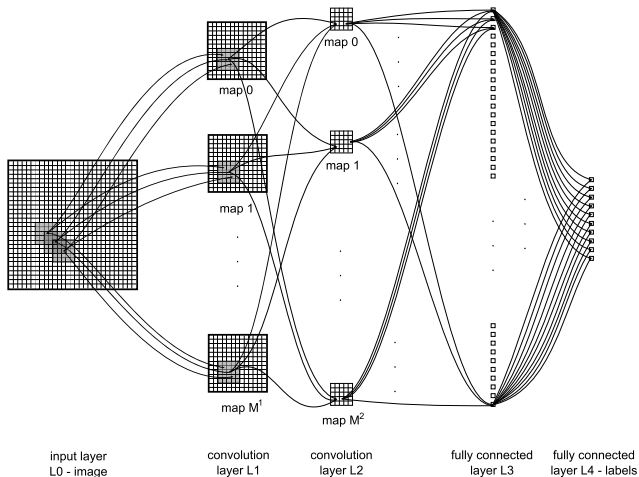
Feedforward neural networks

- A **feedforward neural network** (FNN) is an ANN in which
 - ▶ All connections are directed, and
 - ▶ There are no cycles (i.e., forms a DAG)
- Neurons usually grouped in **layers**
 - ▶ Input neurons: no incoming edges (first layer)
 - ▶ Output neurons: no outgoing edges (last layer)
 - ▶ Hidden neurons: all others (layer = maximum distance from input)
 - ▶ Layers do not need to be fully connected
 - ▶ Traditionally: edges only between subsequent layers (but: edges that skip layers are allowed, too)
- Example: an FNN with one hidden layer (omitting bias inputs)



MNIST, best performer (2011), architecture

Deep convolutional neural network (no preprocessing)



Learning

- Once we settled on an architecture, we need to learn connection strengths (weights)
- Simple approach
 - ▶ Supervised learning with labeled training examples
 - ▶ Use a suitable notion of model performance (e.g., loss function, likelihood)
 - ▶ Learn all weights jointly
 - As before: ERM/RRM, MLE/MAP, Bayesian inference
 - ▶ Most common: empirical/regularized risk minimization
- We will see: simple approach is often “not enough”
 - ▶ High model complexity
 - ▶ Limited (labeled) training data
- Values of hidden units can be thought of as features, but which features are good is unknown and needs to be learned
 - ▶ This makes learning hard
 - ▶ Note: embeddings typically refer to the values of a “certain” hidden layer in a larger FNN (more later)

FNNs, more generally

- Choice of neuron type(s) important
 - ▶ Influences expressability
 - ▶ Influences “learnability”
 - ▶ [Many more types](#) of artificial neurons have been proposed
- Neurons may not follow the template discussed here
 - ▶ E.g., product neuron $y = \prod_i x_i$
 - ▶ E.g., max-pooling neuron $y = \max_i(x_i)$
- Layers may be more general
 - ▶ I.e., any parameterized function from \mathbb{R}^n to \mathbb{R}^m
 - ▶ May compute multiple outputs jointly (e.g., softmax layer)
 - ▶ May have additional internal structure (e.g., Transformer layers)
- Will see: generally, FNNs represented as a “compute graph”